

HUMAIN Create Studio

PREPARED BY  AAVYA

High-Level Architecture

Agentic, Bilingual Content and Experience Platform


Document type	High-Level Architecture
Prepared for	HUMAIN
Region	Kingdom of Saudi Arabia (KSA) — in-region
Classification	Confidential — Client Presentation
Version	1.0
Date	June 2026

1. Overview

HUMAIN Create Studio is an **agentic, AI-native content platform**. A team of specialised AI agents, coordinated by an orchestrator, turns natural-language intent into interactive, on-brand, bilingual outcomes — websites, content, brand guidelines, themes, images and video — all governed by role-based permissions, editorial gates and a complete audit trail, hosted in-Kingdom.


This document describes the system at a high level: context, components, the agentic core, data and integration architecture, request flow, deployment, and the non-functional posture.

2. System Context



The platform is the system-of-record for content; identity, models, render and storage are integrated services.

3. High-Level Architecture




4. Component View

Component	Responsibility	Notes
Edge / Gateway	TLS termination, routing, auth, RBAC/ABAC	In-Kingdom; HTTPS only
Create Studio (console)	Conversational UI + Brand and Design surfaces	Streams agent output (SSE)
Agentic AI Core (AI service)	Orchestrator + specialists + governed tools	Tool-use loop on Claude
Structured CMS	Schema, RBAC/ABAC, editorial workflow, versions, audit	System-of-record
Web Delivery	Fast, cached, RTL-correct public rendering	Bilingual EN/AR
Content Store	Source of truth + vector embeddings	Postgres + pgvector
Search	Arabic-aware + semantic retrieval	—
Object Store	Media and rendered assets	S3-compatible, in-region
Cache / Queue	Performance + background jobs	Redis
Model and Render	LLM completion + image/video rendering	Provider-abstracted


5. Agentic Core

The intelligence is an **orchestrator agent** running a multi-step tool-use loop. It plans, grounds itself in HUMAIN content and brand, delegates to specialist agents, and builds interactive outcomes — never a raw dump. Every tool call is audited; agents propose drafts, humans publish.



- **Orchestrator** runs on the strongest model (Opus) for reasoning/routing.
- **Specialists and generation** run on a fast model (Haiku) for cost efficiency.
- **Governed tools** are the only path to data — read (search/retrieve) and propose (drafts); never direct, unscoped writes.

6. Request Flow — "Build a website"




The same pattern serves every capability — only the build step differs (content, video, image, brand, theme).

7. Data Architecture

- **Content Store (Postgres):** the source of truth — collections, drafts, versions, workflow state, audit log.
- **Vector Index (pgvector):** sovereign, local embeddings (multilingual, 384-dim) power semantic, Arabic-aware retrieval (RAG) — no external embedding calls.
- **Search:** Arabic-aware + semantic retrieval over indexed content.
- **Object Store (S3-compatible, in-region):** media and rendered assets.
- **Cache / Queue (Redis):** performance and reliable background processing (indexing, render polling).

All content and embeddings remain **in-Kingdom**.


8. Integration Architecture



- **Provider-abstracted models** — HUMAIN is never locked to a single vendor.
- **MCP tool surface** — approved external agents act through the same governed, audited tools.
- **Signed webhooks** notify downstream systems on publish/content events.

9. Deployment Architecture

A self-contained, containerised stack, in-Kingdom, behind a TLS edge.



- **Stateless app tier** (console, CMS, AI service) — horizontally scalable.
- **Managed data tier** (Postgres+pgvector, Redis) with health checks.
- **One image, many services** — each container runs a different workload from a shared build.
- **Repeatable deploys** with image cache-busting per release and forced container recreation.

10. Security, Governance and Sovereignty

- **In-Kingdom** hosting and processing — data residency under HUMAIN's control.
- **RBAC + ABAC** at every layer (role, site-scope, department, locale).
- **Editorial gating** — Draft → In Review → Approved; agents propose drafts, humans publish.
- **Immutable audit log** of all content *and* agent actions.
- **Secrets isolation** — credentials/keys separated from content and code.
- **Secure transport** — HTTPS/TLS everywhere; secure session cookies; CORS/CSRF protection.
- **Human-in-the-loop** — nothing reaches the public without approval.

11. Technology Stack

Layer	Technology
Console / Web	Next.js (App Router), React, SSE streaming, EN/AR RTL
CMS	Payload (headless), RBAC/ABAC, versioned drafts, audit
AI service	Node/Fastify, Anthropic Claude (Opus + Haiku), tool-use loop
Agents / tools	Orchestrator + specialists; MCP governed tool surface
Embeddings	Local multilingual model (sovereign, 384-dim)
Data	Postgres + pgvector; Arabic-aware search; Redis
Render	Image/Video via provider-abstracted services
Storage	S3-compatible object storage, in-region
Edge / Infra	Nginx + TLS; containerised; one image, many services

12. Scalability, Reliability and Observability

- **Scalability:** stateless app tier scales horizontally; queue-backed background work absorbs spikes; cost-tuned model routing (Opus for reasoning, Haiku for generation).
- **Reliability:** health-checked data tier; restart-on-failure services; graceful degradation (a missing model/render key surfaces a friendly message, never a crash).
- **Observability:** centralised logs/metrics; every agent run and content action recorded for full traceability.

Appendix A — Payload CMS v3 Security Posture

The platform's content backend is **Payload CMS v3 (3.85.0)**. Payload v3 is a secure, production-grade, actively-maintained headless CMS (Figma-backed) on the Next.js App Router; security is a shared responsibility — Payload provides the primitives, and this deployment configures and hardens them.

A.1 Built-in security controls (Payload v3)

Control	Detail
Authentication	bcrypt-hashed passwords · JWT in httpOnly cookies · API keys · email verification · password reset · account lockout (<code>maxLoginAttempts</code> / <code>lockTime</code>)
Authorization	Collection- and field-level access-control functions → real RBAC + ABAC
Injection defence	Parameterised queries via Drizzle ORM (Postgres)
Web security	CORS + CSRF controls, configurable origins; inherits the Next.js security model
Data integrity	Versioned drafts, editorial workflow state, audit hooks
Maintenance	Actively maintained; pinned, recent 3.85.0

A.2 Applied controls in this platform

Area	Configuration
Transport	HTTPS/TLS · <code>COOKIE_SECURE=true</code> · secure session cookies
Access control	RBAC (6 roles) + ABAC (site, department, locale) at collection and field level
Secrets	<code>PAYLOAD_SECRET</code> and provider keys held in environment, not committed to source
Governance	Editorial gating (draft → review → publish) · immutable audit log
Agentic safety	Agents propose drafts ; a human approves before publish
Sovereignty	Hosted in-Kingdom ; secrets isolated from content and code

A.3 Hardening checklist

- Keep Payload **patched** and monitor its security advisories.
- Strong, rotated `PAYLOAD_SECRET` ; restrict `/admin` exposure; MFA for maintainers.
- **Dependency / SCA scanning** (Dependabot / Snyk / `npm audit`) in CI.
- Replace the first-boot dev schema-push with **generated Payload migrations**.
- Edge **rate-limiting / WAF**; least-privilege database and registry credentials.

No framework is free of vulnerabilities; the correct posture is to **stay current and monitor advisories**. Most real-world risk lies in weak secrets, over-permissive access rules, leaked environment variables or unpatched dependencies — all addressed above.

